

Features and Documentation

This section briefly discusses some of the TCP related features that I've implemented, and how various functions work.

Overall Design

TCP Client basically functions as follows:

1. Given a set of command line arguments in `args`
2. Construct a `TCP_CLIENT` instance which constructs a `UDP_CLIENT` under the hood, and initializes parameters such as `port` number to send to, `port` number to receive from, etc
3. Start a thread for doing blocking receive in a loop. In particular, it runs the following method

```
1 def __receive(client:TCP_CLIENT):
2     while client.state == TCP_CLIENT.ESTABLISHED:
3         # check again
4         client.rcv_lock.acquire()
5         if client.state == TCP_CLIENT.ESTABLISHED:
6             received = client.receive() # blocking
7             logging.info(f'state {client.state}')
8             if client.state != TCP_CLIENT.ESTABLISHED:
9                 client.update_fin_packets(received)
10            logging.info(f'thread received: {received}')
11
12            client.rcv_lock.release()
13            time.sleep(0.1)
14    return
```

here, a lock is required when calling receive, because during the scenario when it starts `FIN` sequences to end the connection, the **main thread** who is constantly sending packets will send a `FIN` packet and **wait for ACK** and etc. Without this lock, there is a chance that the `ACK` packet will be stolen by this thread before the main thread calls `rcv` in the following coed (for instance)

```
1 def __wait_server_ack(self, fin_packet:Packet):
2     logging.debug(f'at __wait_server_ack')
3     fin_seq = fin_packet.header.seq_num
4     self.__state = TCP_CLIENT.FIN_WAIT_1
5
6     # find FIN ACK packet
7     check_threading = True
8     while self.__state == TCP_CLIENT.FIN_WAIT_1:
9         # wait
10        time.sleep(1)
11        self.rcv_lock.acquire()
12        """
13        Obtain LOCK here so that:
14        Case 1. the other thread in tcpclient.py obtained the lock and went rcv.
15            - self.__thread_fin_packets is 100% updated. This works
16        Case 2. I got the lock
17            - the other thread obviously went rcv. Also works
18        """
19        # check list first
20        if check_threading:
21            for packet in self.__thread_fin_packets:
22                # check if is the ACK for fin
23                logging.debug(f'fin ack wait: checking {packet.header}')
24                if packet.header.ack_num >= fin_seq + 1 and packet.header.is_ack():
25                    self.__state = TCP_CLIENT.FIN_WAIT_2
26                    self.__thread_fin_packets.remove(packet)
27                    self.__window = []
```

```

28         self.rcv_lock.release()
29         return packet
30         check_threading = False
31     # receive
32     packet = self.receive()
33     self.rcv_lock.release()
34
35     logging.debug(f'fin ack wait: {packet.header}')
36     if packet.header.ack_num == fin_seq + 1 and packet.header.is_ack():
37         self.__state = TCP_CLIENT.FIN_WAIT_2
38         return packet
39     return

```

4. Start a loop that constantly reads `globals.MSS` (which is set to 512) bytes from the file, and attempt to send:

```

1  # inside the send_file method
2  with open(args.file, 'rb') as openfile:
3      receiv_thread.start()
4      data = openfile.read(globals.MSS)
5      while data != b'':
6          ret = client.send(data)
7          while ret == -1:
8              time.sleep(1)
9              ret = client.send(data)
10         data = openfile.read(globals.MSS)
11     client.terminate()
12     receiv_thread.join()

```

where the `TCP_CLIENT` could reject packet by returning `0` from `client.send(data)`, which will happen when the window size is full. (To see how I implemented Pipelined Sending, checkout the next section.)

5. Once all files contents are sent, terminate the connection by doing the `FIN` handshakes inside the method `terminate`, which will be gone into details in the next section as well.

TCP Receiver/Server does the following:

1. Given a set of command line arguments in `args`
2. Construct a `TCP_SERVER` instance which constructs a `UDP_SERVER` under the hood, and initializes parameters such as `port` number to send to, `port` number to receive from, etc
3. Start the server with `start()`, which will constantly attempt to `service_client` by receiving from the underlying buffer:

```

1  def start(self, args):
2      """Start the server
3      Start to listen and accept clients.
4      Reset when client initiates FIN requests and completed the handshake
5
6      Args:
7          args (namespace): command line arguments for the program, e.g. which file to write
8          to
9      """
10     server = self._socket
11     server.bind(self._serveraddress)
12     # other application related init
13     init(args)
14     self.__state = TCP_SERVER.LISTEN
15     print("The server is ready to receive")
16
17     while True:
18         try:
19             self.__state = TCP_SERVER.ESTABLISHED
20             service_client(self, args)
21             self.__state = TCP_SERVER.LISTEN

```

```

21         except Exception as err:
22             print(err)
23             pass
24     return

```

where the `service_client` function basically is the entry point for all the receiving and acking:

```

1  def service_client(server:TCP_SERVER, args):
2      """Specifies what to do when received something from client
3
4      Essentially does 1) receive 2) check packet received
5      3) write to file 4) send ACK
6
7      Args:
8          server (TCP_SERVER): running instance of TCP_SERVER
9          args (namespace): command line arguments
10     """
11     # receive packet
12     received, client_address = server.receive()
13     logging.info(f"[LOG] serviced {client_address}")
14     logging.info(f"{received or 'Discarded or Residual'}")
15
16     if received is not None:
17         # write to file
18         to_file(received, dst=args.file)
19
20     # send ACK
21     if server.state == TCP_SERVER.ESTABLISHED:
22         server.send('')
23     return

```

which basically a) receives from the buffer, b) write to file if non-corrupt (not-None) packet received, c) send `ACK`

- When a client initiated a `FIN` sequence, the `server.receive()` will detect a `FIN` packet and basically does an `ACK` back and sends a `FIN`. Implementation details of this will be gone over in the next section.

Once this "handshake" is done, it will reset its current status of `seq_num` and `ack_num` and etc, such that it can be ready to service the next "new client".

- This `service_client` mentioned above will continue running forever, and should work across multiple runs of `tcpsender.py` without the need to restarting the server everytime.

TCP Related Features

- TCP Packet**

This is implemented as the following class

```

1
2  class Packet(util.Comparable):
3      """TCP Packet
4
5      This abstraction gives you a human-readable packet on the "surface",
6      but during transmission, it will be "serialized" by struct.pack to become
7      bytes.
8      """
9
10     def __init__(self, header:TCPHeader, payload:str) → None:
11         """Construct a packet from header and payload
12
13         Args:
14             header (TCPHeader): a constructed TCP header
15             payload (str or bytes): payload
16         """
17         self.__header = header
18         self.__payload = payload
19

```

```

20     @property
21     def header(self):
22         return self.__header
23
24     @property
25     def payload(self):
26         return self.__payload
27
28     @header.setter
29     def header(self, value):
30         self.__header = value
31
32     @payload.setter
33     def payload(self, value):
34         self.__payload = value
35
36     def compute_checksum(self):
37         checksum = self.__compute_checksum()
38         checksum = ~checksum & 0xffff # 1s complement and mask
39         self.__header.set_checksum(checksum)
40         return
41
42     def __compute_checksum(self):
43         """Computes the 1s complement treating self.__header=0
44
45         Returns:
46             [int]: 1s complement of checksummed packet
47         """
48         prev_checksum = self.__header.checksum
49         self.__header.set_checksum(value=0)
50         # computes checksum without header
51         all_bytes = serialize(self)
52         checksum = 0
53         for i in range(0, len(all_bytes), 2):
54             if i + 1 == len(all_bytes):
55                 checksum += all_bytes[i]
56                 break
57             checksum += (all_bytes[i] << 8 + all_bytes[i+1])
58         # reset
59         self.__header.set_checksum(prev_checksum)
60         return checksum
61
62     def is_corrupt(self):
63         current_checksum = self.__compute_checksum()
64         logging.debug(f'checksum result {current_checksum & self.__header.checksum}')
65         return current_checksum & self.__header.checksum != 0
66
67     def __str__(self):
68         content = f"""
69         ---
70         [HEADER]: {self.__header}
71         [Payload]: {self.__payload}
72         ---"""
73         return content

```

basically it contains human-readable formatting of a Packet, which is very useful for logging and debugging. When actually sending/receiving the packet, it will be serialized into bytes by the following two methods:

```

1  def serialize(packet:Packet):
2      """Converts the Human-readable Packet to bytes
3
4      Args:
5          packet (Packet): Packet abstraction
6
7      Returns:
8          bytes: actual bytes of the packet
9          (i.e. 20 bytes header + up to 512 byte payload)
10     """

```

```

11     line_1 = struct.pack('HH', packet.header.src_port, packet.header.dst_port)
12     line_2 = struct.pack('I', packet.header.seq_num)
13     line_3 = struct.pack('I', packet.header.ack_num)
14     # convert flag
15     flag = packet.header.flags
16     flag_map = int(f"{flag.ack}{flag.cwr}{flag.ece}{flag.fin}{flag.syn}", 2)
17     line_4 = struct.pack('BBH', packet.header.header_len, flag_map, packet.header.rcvwd)
18     line_5 = struct.pack('HH', packet.header.checksum, 0)
19     # final
20     final = line_1 + line_2 + line_3 + line_4 + line_5
21     if len(packet.payload) != 0:
22         final += packet.payload
23     return final
24
25 def deserialize(packet):
26     """Converts bytes to a HUMAN-readable Packet
27
28     Args:
29         packet (bytes): network transmitted bytes
30
31     Returns:
32         Packet: human-readable Packet
33     """
34     total_size = len(packet)
35     src_port, dst_port, \
36         seq_num, ack_num, \
37         header_len, flags, rcvwd, \
38         checksum, urg, \
39         data = struct.unpack(f'HHIIBBHHH{total_size-20}s', packet)
40     flags = format(flags, '#07b')
41     header = TCPHeader(
42         src_port=src_port,
43         dst_port=dst_port,
44         seq_num=seq_num,
45         ack_num=ack_num,
46         _flags=Flags(
47             cwr=int(flags[3]),
48             ece=int(flags[4]),
49             ack=int(flags[2]),
50             syn=int(flags[6]),
51             fin=int(flags[5])),
52         rcvwd=rcvwd)
53     header.set_checksum(checksum)
54     packet = Packet(header, data)
55     return packet

```

the function `serialize` converts a `Packet` instance to bytes using `struct.pack`, and the other one unpacks the bytes into a `Packet` by `struct.unpack`.

The serialize/deserialize happens **only right before** the `send_to` and **right after** the `recv_from` of the UDP channel, which means the entire TCP code can treat all data as a `Packet` (which makes the program easier):

```

1  # insude UDP_SERVER
2  def send_packet(self, packet:Packet, client_address):
3      socket = self.__socket
4      packet = structure.packet.serialize(packet)
5      ret = socket.sendto(packet, client_address)
6      return ret
7  def receive_packet(self):
8      raw_packet, _ = self.__socket.recvfrom(self.__buffersize)
9      return structure.packet.deserialize(raw_packet)

```

and similarly in the `UDP_CLIENT`:

```

1  def send_packet(self, packet:Packet, client_address):
2      socket = self.__socket

```

```

3         packet = structure.packet.serialize(packet)
4         ret = socket.sendto(packet, client_address)
5         return ret
6
7     def receive_packet(self):
8         server = self.__socket
9         raw_packet, client_address = server.recvfrom(self.__buffer_size)
10        try:
11            # e.g. corruption
12            packet = structure.packet.deserialize(raw_packet)
13            logging.debug(f'rcvd {packet}')
14        except:
15            packet = None
16        return packet, client_address

```

- **Timer** (used for timeouts)

This in TCP is basically based on the `threading.Timer` module:

```

1 class TCPTimer(object):
2     """TCP timer implementation. Used for multithreading mainly
3     """
4
5     def __init__(self, interval: float, function: Callable[..., Any], *args, **kwargs) → None:
6         """TCP Timer implementation. Essentially triggers @function when timedout.
7
8         Args:
9             interval (float): TimeoutInterval
10            function (Callable[..., Any]): function to call when timedout
11        """
12        self.__interval = interval
13        self.__function = function
14        self.__args = args
15        self.__kwargs = kwargs
16        self.__timer = None
17
18    @property
19    def interval(self):
20        return self.__interval
21
22    def start(self):
23        if self.__timer is not None and self.__timer.is_alive():
24            logging.error('timer already running')
25            return
26        self.__timer = Timer(self.__interval, self.__function, args=self.__args,
27                             kwargs=self.__kwargs)
28        self.__timer.start()
29        logging.debug('timer started')
30        return
31
32    def is_alive(self):
33        if self.__timer is None:
34            return False
35        return self.__timer.is_alive()
36
37    def cancel(self):
38        if self.__timer is None:
39            return
40        self.__timer.cancel()
41        return
42
43    def restart(self, new_interval=None):
44        if self.__timer is not None and self.__timer.is_alive():
45            self.__timer.cancel()
46
47        # calling self.start() causes problem as the thread might NOT be finished
48
49        # (e.g. function still executing)

```

```

50     self.__interval = new_interval or self.__interval
51     self.__timer = Timer(self.__interval, self.__function, args=self.__args,
        kwargs=self.__kwargs)
52     self.__timer.start()
53     return

```

where whenever you start a `TCPTimer`, it basically starts another thread using `Timer`. When that `Timer` object timed out, it will essentially call the `self.__function` (which will be the `retransmit` function). Therefore, in this way, retransmissions of packets will not interfere with the main program of sending/receiving.

Inside the TCP client, some of the usages look like:

```

1  class TCP_CLIENT(UDP_CLIENT):
2      def __init__(self, udpl_ip, udpl_port, window_size, ack_lstn_port):
3          """TCP reliable sender implementation
4
5          Args:
6              udpl_ip (str): udpl IP address to send to (proxy address)
7              udpl_port (int): udpl port address to send to
8              window_size (int): number of packets allowed in current window
9              ack_lstn_port (int): port number of receiving ACK from server
10         """
11         super().__init__(udpl_ip, udpl_port, ack_lstn_port)
12         self.__seq_num = 0
13         self.__ack_num = 0 # assumes both sides start with seq=0
14         self.__timer = timer.TCPTimer(TCP_CLIENT.INIT_TIMEOUT_INTERVAL, self.retransmit)
15         # other initialization omitted
16
17     # other class methods omitted
18     def __post_send(self, packet:Packet):
19         logging.debug("at __post_send")
20         # some code omitted
21
22         # 3. check if timer is running
23         self.__rtt_sampling.double_interval(enabled=False, restore=False)
24         if not self.__timer.is_alive():
25             self.__timer.restart(new_interval=self.__rtt_sampling.get_interval())
26         # some code omitted
27     return

```

where notice that when there is a timeout, it will go doubling the interval by `self.__rtt_sampling.double_interval(enabled=False, restore=False)` and then the timer will restart (if not running) with an interval of `self.__timer.restart(new_interval=self.__rtt_sampling.get_interval())`.

To see how the RTT Sampler works, see the next bullet point on RTT Sampler.

- **RTT Sampler**

Again, this is implemented as an object so that detailed updating mechanism can be hidden away from the main logics of TCP sending. In details, it is implemented as follows:

```

1  class RTTSampler(object):
2      """TCP RTT Sampler
3
4      This class essentially allows you to input a measured RTT and updates
5      TimeoutInterval internally. So the next time, you can get the computed
6      TimeoutInterval by :func:self.get_interval
7      """
8      def __init__(self, init_interval) → None:
9          super().__init__()
10         self.__timeout_interval = init_interval
11
12         # used for estimation
13         self.__estimated_rtt = init_interval
14         self.__alpha = 0.125
15         self.__dev_rtt = 0
16         self.__beta = 0.25
17         self.__gamma = 2

```

```

18
19     # used for doubling timeout
20     self.__within_timeout = False
21     pass
22
23     def double_interval(self, enabled=True, restore=True):
24         self.__within_timeout = enabled
25         if enabled is False and restore: # when sending new packets, do not restore
26             self.__timeout_interval = round(self.__estimated_rtt + self.__gamma *
self.__dev_rtt, 3)
27         return
28
29     def update_interval(self, sample_rtt):
30         # we received something, switch back to using normal timeout
31         self.__within_timeout = False
32
33         alpha = self.__alpha
34         beta = self.__beta
35         self.__estimated_rtt = (1-alpha) * self.__estimated_rtt + alpha * sample_rtt
36         self.__dev_rtt = (1-beta) * self.__dev_rtt + beta * (abs(sample_rtt -
self.__estimated_rtt))
37         self.__timeout_interval = round(self.__estimated_rtt + self.__gamma * self.__dev_rtt,
3)
38         logging.debug(f"""
39             sample with {sample_rtt}
40             new self.__estimated_rtt {self.__estimated_rtt}
41             new self.__dev_rtt {self.__dev_rtt}
42             rounded new timeout interval {self.__timeout_interval}
43             """)
44         return
45
46     def get_interval(self):
47         if self.__within_timeout:
48             self.__timeout_interval *= 2
49         return self.__timeout_interval

```

where the most important method is basically the `update_interval`. This basically performs the calculation of the new timeout interval, and it is done everytime when `TCP_CLIENT` received an `ACK`, and there is a packet that we are actively tracking:

```

1  # inside TCP_CLIENT
2  def __post_rcv(self, packet:Packet):
3      # some code omitted here
4      # 1. update window, received ACK
5      if packet.header.ack_num > self.__send_base:
6          # 2. new ACK received
7          # some code omitted here
8
9          # 3. update RTT
10         start_time = self.__waiting_packets.get(packet.header.ack_num)
11         if start_time is not None: # not retransmitted
12             end_time = time.time()
13             self.__rtt_sampling.update_interval(end_time - start_time)
14             self.__waiting_packets.pop(packet.header.ack_num, None)
15         return
16     # some code omitted here
17     return

```

where basically it will check the `self.__waiting_packets` dictionary, which is added and updated by the `send` and `retransmit` methods, and see if we should sample this RTT or not.

- **Pipelined Sending** (window)

This is now simple due to the multithreading receive. In essence, whenever we send something, it a) check the current window size to see if it is full, b) if not full, send c) perform `__post_send` actions such as updating the next sequen number, adding packet into window, and start a timing track for RTT sampler


```

1  def send(self, payload:str):
2      """Reliably send a packet with payload @payload
3
4      Args:
5          payload (str or bytes): payload
6
7      Returns:
8          int: success=0
9      """
10     # 0. consult window
11     if len(self.__window) == self.__window_size:
12         return -1
13     # 1. construct packet
14     _, src_port = self.get_info()
15     header = TCPHeader(
16         src_port=src_port,
17         dst_port=self.dst_addr[1],
18         seq_num=self.__seq_num,
19         ack_num=self.__ack_num,
20         _flags=Flags(cwr=0, ece=0, ack=0, syn=0, fin=0),
21         rcvwd=10)
22     packet = Packet(header, payload)
23     packet.compute_checksum()
24
25     # 2. send packet
26     self.send_packet(packet)
27
28     # 3. update seq_num, etc
29     self.__post_send(packet)
30     return 0

```

(notice that we were able to construct a `Packet` in an entirely human-readable form because all the byte conversion is done secretly in the end by `serialize/deserialize` !)

- **Checksum**

To prevent against corrupted packet, the internet checksum basically performs the summing over data and check against the `checksum` field of the packet.

On the sender side, a checksum is computed everytime we constructed a packet

```

1  # code snipped inside send() of TCP_CLIENT
2  _, src_port = self.get_info()
3  header = TCPHeader(
4      src_port=src_port,
5      dst_port=self.dst_addr[1],
6      seq_num=self.__seq_num,
7      ack_num=self.__ack_num,
8      _flags=Flags(cwr=0, ece=0, ack=0, syn=0, fin=0),
9      rcvwd=10)
10 packet = Packet(header, payload)
11 packet.compute_checksum()

```

where the `compute_checksum` basically does

```

1  def compute_checksum(self):
2      checksum = self.__compute_checksum()
3      checksum = ~checksum & 0xffff # 1s complement and mask
4      self.__header.set_checksum(checksum)
5      return

```

where the `__compute_checksum()` basically computes the raw sum without 1s complement:

```

1  def __compute_checksum(self):
2      """Computes the 1s complement treating self.__header=0
3

```

```

4     Returns:
5         [int]: 1s complement of checksummed packet
6     """
7     prev_checksum = self.__header.checksum
8     self.__header.set_checksum(value=0)
9     # computes checksum without header
10    all_bytes = serialize(self)
11    checksum = 0
12    for i in range(0, len(all_bytes), 2):
13        if i + 1 == len(all_bytes):
14            checksum += all_bytes[i]
15            break
16        checksum += (all_bytes[i] << 8 + all_bytes[i+1])
17    # reset
18    self.__header.set_checksum(prev_checksum)
19    return checksum

```

On the receiver/server side, whenever a packet is received, it can invoke the `is_corrupt` method to check the checksum easily:

```

1 def is_corrupt(self):
2     current_checksum = self.__compute_checksum()
3     logging.debug(f'checksum result {current_checksum & self.__header.checksum}')
4     return current_checksum & self.__header.checksum != 0

```

- **FIN**

When a client finishes sending all pieces of a file, it starts calling `terminate` function, which basically starts performing the `FIN` "handshake":

```

1 def terminate(self):
2     """Terminate the connection
3
4     After the FIN handshake, close the underlying UDP socket.
5
6     Returns:
7         None: None
8     """
9     # 0. wait for all other retransmission to be done
10    while len(self.__window) > 0:
11        # the other thread will timeout and retransmit
12        time.sleep(1)
13
14    # change state so that the other thread will not receive packets
15    self.__state = TCP_CLIENT.BEGIN_CLOSE
16
17    # 1. construct FIN packet
18    _, src_port = self.get_info()
19    header = TCPHeader(
20        src_port=src_port,
21        dst_port=self.dst_addr[1],
22        seq_num=self.__seq_num,
23        ack_num=self.__ack_num,
24        _flags=Flags(cwr=0, ece=0, ack=0, syn=0, fin=1),
25        rcvwd=10)
26    packet = Packet(header, b'')
27    packet.compute_checksum()
28    self.__fin_start_seq = self.__seq_num
29
30    # 2. send packet
31    self.send_packet(packet)
32
33    # 3. start timers
34    self.__post_send(packet)
35
36    # 4. wait for acks and etc

```

```
37     self.__post_fin(packet)
38     return super().terminate()
```

where `super().terminate()` terminates the entire socket, so it is the end of transmission.

And similarly, the server will detect the `FIN` handshake during a `receive()` call:

```
1  def receive(self):
2      """Blocking receive a packet
3
4      Returns:
5          (Packet, tuple): returns (Packet, client_address) if packet is not corrupt.
6          Else, returns (None, client_address)
7      """
8      # 1. receive packet
9      packet, client_address = self.receive_packet()
10     # 2. check if packet is corrupt
11     if packet is not None and not packet.is_corrupt():
12         # 3. if not, update ack_num
13         packet = self.__post_rcv(packet)
14     else:
15         packet = None
16     return packet, client_address
```

and inside `__post_rcv()`, it does:

```
1  def __post_rcv(self, packet:Packet):
2      self.__ack_num = self.__next_ack(packet) # position of next byte
3      if packet.header.is_fin() and packet.header.seq_num + 1 >= self.__ack_num:
4          logging.info('closing connection')
5          logging.info(packet)
6          packet = self.close_connection(packet)
7          logging.info('connection closed')
8      return packet
```